

Лабораторная работа №5  
по дисциплине  
«Методы машинного обучения»  
на тему  
«Линейные модели, SVM и деревья решений»

Выполнил:  
студент группы ИУ5-23М  
Наседкин И. А.

---

# 1. Линейные модели, SVM и деревья решений

## 2. Задание:

Выберите набор данных (датасет) для решения задачи классификации или регрессии.  
В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.

С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.

Обучите следующие модели:

одну из линейных моделей;

SVM;

дерево решений.

Оцените качество моделей с помощью трех подходящих для задачи метрик. Сравните качество полученных моделей.

Произведите для каждой модели подбор одного гиперпараметра с использованием `GridSearchCV` и кросс-валидации.

Повторите пункт 4 для найденных оптимальных значений гиперпараметров. Сравните качество полученных моделей с качеством моделей, полученных в пункте 4.

Дополнительные задания:

Проведите эксперименты с важностью признаков в дереве решений.

Визуализируйте дерево решений.

```
[1]: import numpy as np
import pandas as pd
from typing import Dict, Tuple
from scipy import stats
from IPython.display import Image
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score,
    ↪ classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import mean_absolute_error, mean_squared_error,
    ↪ mean_squared_log_error, median_absolute_error, r2_score
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import cross_val_score, cross_validate
from sklearn.svm import SVC, NuSVC, LinearSVC, OneClassSVM, SVR, NuSVR,
    ↪ LinearSVR
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import SGDRegressor
from sklearn.linear_model import SGDClassifier
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
def make_dataframe(ds_function):
    ds = ds_function()
    df = pd.DataFrame(data= np.c_[ds['data'], ds['target']],
```

```

        columns= list(ds['feature_names']) + ['target'])
    return df
from sklearn.datasets import fetch_california_housing
from sklearn.externals.six import StringIO
import graphviz
import pydotplus
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor,
    export_graphviz
from sklearn.model_selection import GridSearchCV

```

```

/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19:
FutureWarning: pandas.util.testing is deprecated. Use the functions in
    the
public API at pandas.testing instead.
    import pandas.util.testing as tm
/usr/local/lib/python3.6/dist-packages/sklearn/externals/six.py:31:
FutureWarning: The module is deprecated in version 0.21 and will be
    removed in
version 0.23 since we've dropped support for Python 2.7. Please rely on
    the
official version of six (https://pypi.org/project/six/).
    "(https://pypi.org/project/six/).", FutureWarning)

```

```
[0]: data = fetch_california_housing()
```

### 3. Датасет данных о жилье в Калифорнии(пропуски отсутствуют, нет категориальных признаков для кодирования):

### 4. Разбиение выборки:

```
[3]: df_california = pd.DataFrame(data.data, columns=data.feature_names)
df_california['target'] = pd.Series(data.target)
df_california.head()
```

```
[3]:   MedInc  HouseAge  AveRooms  AveBedrms  ...  AveOccup  Latitude  Longitude
    target
0   8.3252     41.0   6.984127   1.023810  ...   2.555556     37.88    -122.
    23
4.526
1   8.3014     21.0   6.238137   0.971880  ...   2.109842     37.86    -122.
    22
3.585
2   7.2574     52.0   8.288136   1.073446  ...   2.802260     37.85    -122.
    24
3.521
```

```

3  5.6431      52.0  5.817352  1.073059  ...  2.547945      37.85      -122.
   ↪ 25
3.413
4  3.8462      52.0  6.281853  1.081081  ...  2.181467      37.85      -122.
   ↪ 25
3.422

```


[5 rows x 9 columns]

```

[4]: fig, ax = plt.subplots(figsize=(15,7))
      sns.heatmap(df_california.corr(method='pearson'), ax=ax, annot=True,
      ↪fmt='.2f')

```

[4]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f5cef4d8668>




```

[5]: fig, ax = plt.subplots(figsize=(10,10))
      sns.scatterplot(ax=ax, x='MedInc', y='target', data=df_california)

```

[5]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f5cff5e0ac8>



Между признаком “MedInc” и целевым признаком “target” существует зависимость, близкая к линейной, коэффициент корреляции = 0,69

Попробуем восстановить данную линейную зависимость.

аналитическое вычисление коэффициентов:

```
[0]: def analytic_regr_coef(x_array : np.ndarray,
                             y_array : np.ndarray) -> Tuple[float, float]:
    x_mean = np.mean(x_array)
    y_mean = np.mean(y_array)
    var1 = np.sum([(x-x_mean)**2 for x in x_array])
    cov1 = np.sum([(x-x_mean)*(y-x_mean) for x, y in zip(x_array,
    ↪ y_array)])
    b1 = cov1 / var1
    b0 = y_mean - b1*x_mean
    return b0, b1
```

```
[0]: x_array = df_california['MedInc'].values
      y_array = df_california['target'].values
```


```
[8]: b0, b1 = analytic_regr_coef(x_array, y_array)
      b0, b1
```

```
[8]: (0.45085576703268027, 0.41793849201896244)
```

```
[0]: # Вычисление значений y на основе x для регрессии
      def y_regr(x_array : np.ndarray, b0: float, b1: float) -> np.ndarray:
          res = [b1*x+b0 for x in x_array]
          return res
```

```
[0]: y_array_regr = y_regr(x_array, b0, b1)
```

```
[11]: plt.plot(x_array, y_array, 'g.')
      plt.plot(x_array, y_array_regr, 'b', linewidth=2.0)
      plt.show()
```




```
[12]: # Синими отрезками показаны ошибки между
      # истинными и предсказанными значениями
      K_mnk=10

      plt.plot(x_array[1:K_mnk+1], y_array[1:K_mnk+1], 'bo')
      plt.plot(x_array[1:K_mnk+1], y_array_regr[1:K_mnk+1], '-ro',
               ↪ linewidth=2.0)

      for i in range(len(x_array[1:K_mnk+1])):
          x1 = x_array[1:K_mnk+1][i]
          y1 = y_array[1:K_mnk+1][i]
          y2 = y_array_regr[1:K_mnk+1][i]
          plt.plot([x1,x1],[y1,y2], 'b-')

      plt.show()
```



```
[0]: # Простейшая реализация градиентного спуска
def gradient_descent(x_array : np.ndarray,
                    y_array : np.ndarray,
                    b0_0 : float,
                    b1_0 : float,
                    epochs : int,
                    learning_rate : float = 0.001
                    ) -> Tuple[float, float]:
    # Значения для коэффициентов по умолчанию
    b0, b1 = b0_0, b1_0
    k = float(len(x_array))
    for i in range(epochs):
        # Вычисление новых предсказанных значений
        # используется векторизованное умножение и сложение для вектора
        ↪и константы
        y_pred = b1 * x_array + b0
        # Расчет градиентов
        # np.multiply - поэлементное умножение векторов
        dL_db1 = (-2/k) * np.sum(np.multiply(x_array, (y_array -
        ↪y_pred)))
        dL_db0 = (-2/k) * np.sum(y_array - y_pred)
        # Изменение значений коэффициентов:
        b1 = b1 - learning_rate * dL_db1
        b0 = b0 - learning_rate * dL_db0
        # Результирующие значения
        y_pred = b1 * x_array + b0
    return b0, b1, y_pred
```

```
[0]: def show_gradient_descent(epochs, b0_0, b1_0):
    grad_b0, grad_b1, grad_y_pred = gradient_descent(x_array, y_array,
    ↪b0_0, b1_0, epochs)
    print('b0 = {} - (теоретический), {} - (градиентный спуск)'.
    ↪format(b0, grad_b0))
```


```

print('b1 = {} - (теоретический), {} - (градиентный спуск)'.
→format(b1, grad_b1))
print('MSE = {}'.format(mean_squared_error(y_array_regr,
→grad_y_pred)))
plt.plot(x_array, y_array, 'g.')
plt.plot(x_array, y_array_regr, 'b', linewidth=2.0)
plt.plot(x_array, grad_y_pred, 'r', linewidth=2.0)
plt.show()

```

[15]: *# Примеры использования градиентного спуска*  
show\_gradient\_descent(10, 0, 0)


b0 = 0.45085576703268027 - (теоретический), 0.035046884360901316 -  
→(градиентный  
спуск)  
b1 = 0.41793849201896244 - (теоретический), 0.16024063837885774 -  
→(градиентный  
спуск)  
MSE = 2.237016080434203



[16]: show\_gradient\_descent(1000, 0, 0)


b0 = 0.45085576703268027 - (теоретический), 0.21125447583768953 -  
→(градиентный  
спуск)  
b1 = 0.41793849201896244 - (теоретический), 0.46832724481210447 -  
→(градиентный  
спуск)  
MSE = 0.011149582877403865





```
[17]: show_gradient_descent(50000, 0, 0)
```


$b_0 = 0.45085576703268027$  - (теоретический),  $0.4508557641340037$  -  
 → (градиентный спуск)  
 $b_1 = 0.41793849201896244$  - (теоретический),  $0.4179384926285615$  -  
 → (градиентный спуск)  
 $MSE = 1.6318485911138797e-18$



```
[18]: show_gradient_descent(100000, 0, 0)
```


$b_0 = 0.45085576703268027$  - (теоретический),  $0.45085576703259417$  -  
 → (градиентный спуск)  
 $b_1 = 0.41793849201896244$  - (теоретический),  $0.41793849201898114$  -  
 → (градиентный спуск)

спуск)  
MSE = 1.4521197259897182e-27




```
[19]: show_gradient_descent(1000000, 0, 0)
```

b0 = 0.45085576703268027 - (теоретический), 0.45085576703259417 -  
→(градиентный  
спуск)  
b1 = 0.41793849201896244 - (теоретический), 0.41793849201898114 -  
→(градиентный  
спуск)  
MSE = 1.4521197259897182e-27




```
[20]: show_gradient_descent(1000, -30, 5)
```

$b_0 = 0.45085576703268027$  - (теоретический),  $-20.286696565576065$  -  
 → (градиентный спуск)  
 $b_1 = 0.41793849201896244$  - (теоретический),  $4.779097786324618$  -  
 → (градиентный спуск)  
 $MSE = 83.52092622146682$



[21]: `show_gradient_descent(100000, -30, 5)`

$b_0 = 0.45085576703268027$  - (теоретический),  $0.45085576703259417$  -  
 → (градиентный спуск)  
 $b_1 = 0.41793849201896244$  - (теоретический),  $0.41793849201898114$  -  
 → (градиентный спуск)  
 $MSE = 1.4521197259897182e-27$



## 5. Линейная регрессия при помощи библиотеки scikit-learn

```
[0]: california_x_columns = ['MedInc',
    'HouseAge',
    'AveRooms',
    'AveBedrms',
    'Population',
    'AveOccup',
    'Latitude',
    'Longitude']

[23]: # Обучим линейную регрессию и сравним коэффициенты с рассчитанными ранее
reg1 = LinearRegression().fit(x_array.reshape(-1, 1), y_array.
    ↪ reshape(-1, 1))
(b1, reg1.coef_), (b0, reg1.intercept_)

[23]: ((0.41793849201896244, array([[0.41793849]])),
    (0.45085576703268027, array([0.45085577])))

[0]: sc1 = MinMaxScaler()
sc1_data = sc1.fit_transform(df_california[california_x_columns])

[0]: scoring = {'MAE': 'neg_mean_absolute_error',
    'MSE': 'neg_mean_squared_error',
    'R2': 'r2'}

[26]: scores_3 = cross_validate(LinearRegression(),
    sc1_data, df_california['target'],
    ↪ scoring=scoring,
    cv=3, return_train_score=True)
scores_3

[26]: {'fit_time': array([0.00827169, 0.00720763, 0.00623298]),
    'score_time': array([0.00176811, 0.00223017, 0.00171518]),
    'test_MAE': array([-0.57887213, -0.50877351, -0.56305949]),
    'test_MSE': array([-0.59988822, -0.49188049, -0.59155481]),
    'test_R2': array([0.55502126, 0.58837838, 0.58544641]),
    'train_MAE': array([-0.51394065, -0.54832542, -0.52079949]),
    'train_MSE': array([-0.49933994, -0.54483711, -0.50929041]),
    'train_R2': array([0.62059396, 0.61067588, 0.59991471])}
```

## 6. Данные метрик линейной регрессии для сравнения с другими моделями

```
[27]: -np.mean(scores_3['train_MAE']), -np.mean(scores_3['test_MAE']), \
    -np.mean(scores_3['train_MSE']), -np.mean(scores_3['test_MSE']), \
    np.mean(scores_3['train_R2']), np.mean(scores_3['test_R2'])
```

```
[27]: (0.5276885217738251,  
      0.5502350422781422,  
      0.5178224843539592,  
      0.5611078368165351,  
      0.6103948517000314,  
      0.5762820158960761)
```

```
[28]: # Для небольшой выборки качество обучения сильно уступает  
      ↪ нестохастическому градиентному спуску.  
      print('Размер выборки - {}'.format(x_array.shape[0]))  
      reg2 = SGDRegressor().fit(x_array.reshape(-1, 1), y_array)  
      (b1, reg2.coef_), (b0, reg2.intercept_)
```

Размер выборки - 20640

```
[28]: ((0.41793849201896244, array([0.42488581])),  
      (0.45085576703268027, array([0.45235938])))
```

```
[0]: from sklearn.linear_model import Lasso
```

```
[30]: reg3 = Lasso().fit(x_array.reshape(-1, 1), y_array)  
      (b1, reg3.coef_), (b0, reg3.intercept_)
```

```
[30]: ((0.41793849201896244, array([0.14086477])),  
      (0.45085576703268027, 1.5233170059994505))
```

```
[0]: from sklearn.linear_model import Ridge
```

```
[32]: reg4 = Ridge().fit(x_array.reshape(-1, 1), y_array)  
      (b1, reg4.coef_), (b0, reg4.intercept_)
```

```
[32]: ((0.41793849201896244, array([0.41793288])),  
      (0.45085576703268027, 0.4508774829636899))
```

```
[0]: from sklearn.linear_model import ElasticNet
```

```
[34]: reg5 = ElasticNet().fit(x_array.reshape(-1, 1), y_array)  
      (b1, reg5.coef_), (b0, reg5.intercept_)
```

```
[34]: ((0.41793849201896244, array([0.24540411])),  
      (0.45085576703268027, 1.1186795859141516))
```


## 7. Машина опорных векторов

### 8. Реализация в коде

```
[0]: california_x = df_california['MedInc'].values  
california_y = df_california['target'].values
```

```
[36]: fig, ax = plt.subplots(figsize=(5,5))  
sns.scatterplot(ax=ax, x=california_x, y=california_y)
```


```
[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5ceaf35438>
```



```
[0]: def plot_regr(clf):  
    title = clf.__repr__  
    clf.fit(california_x.reshape(-1, 1), california_y)  
    california_y_pred = clf.predict(california_x.reshape(-1, 1))  
    fig, ax = plt.subplots(figsize=(5,5))  
    ax.set_title(title)  
    ax.plot(california_x, california_y, 'b.')  
    ax.plot(california_x, california_y_pred, 'ro')  
    plt.show()
```

```
[38]: plot_regr(LinearSVR(C=1.0, max_iter=10000))
```

```
<bound method BaseEstimator.__repr__ of LinearSVR(C=1.0, dual=True, epsilon=0.0, fit_intercept=True,
intercept_scaling=1.0, loss='epsilon_insensitive', max_iter=10000,
random_state=None, tol=0.0001, verbose=0)>
```



```
[0]: #plot_regr(LinearSVR(C=1.0, loss='squared_epsilon_insensitive',
↪max_iter=10000))
```

```
[0]: #plot_regr(SVR(kernel='linear', C=1.0))
```

```
[0]: #plot_regr(SVR(kernel='rbf', gamma=0.2, C=1.0))
```

```
[0]: #plot_regr(SVR(kernel='rbf', gamma=0.8, C=1.0))
```

```
[0]: #plot_regr(NuSVR(kernel='rbf', gamma=0.8, nu=0.1, C=1.0))
```

```
[0]: #plot_regr(NuSVR(kernel='rbf', gamma=0.8, nu=0.9, C=1.0))
```

```
[0]: #plot_regr(SVR(kernel='poly', degree=2, gamma='auto', C=1.0))
```

```
[0]: #plot_regr(SVR(kernel='poly', degree=3, gamma=0.2, C=1.0))
```

```
[0]: #plot_regr(SVR(kernel='poly', degree=4, gamma=0.2, C=1.0))
```

```
[48]: df_california.head()
```

```
[48]:   MedInc  HouseAge  AveRooms  AveBedrms  ...  AveOccup  Latitude
↪Longitude
target
0  8.3252    41.0    6.984127    1.023810  ...    2.555556    37.88    -122.
↪23
4.526
1  8.3014    21.0    6.238137    0.971880  ...    2.109842    37.86    -122.
↪22
3.585
```

```

2  7.2574      52.0  8.288136  1.073446 ... 2.802260      37.85      -122.
  ↪ 24
3.521
3  5.6431      52.0  5.817352  1.073059 ... 2.547945      37.85      -122.
  ↪ 25
3.413
4  3.8462      52.0  6.281853  1.081081 ... 2.181467      37.85      -122.
  ↪ 25
3.422

[5 rows x 9 columns]

```

```

[0]: california_x_columns = ['MedInc',
    'HouseAge',
    'AveRooms',
    'AveBedrms',
    'Population',
    'AveOccup',
    'Latitude',
    'Longitude']

```

```

[50]: # Диапазоны значений достаточно сильно различаются
df_california[california_x_columns].describe()

```

```

[50]:
count    20640.000000    20640.000000 ... 20640.000000    20640.000000
mean       3.870671      28.639486 ... 35.631861    -119.569704
std        1.899822      12.585558 ... 2.135952      2.003532
min         0.499900       1.000000 ... 32.540000    -124.350000
25%         2.563400      18.000000 ... 33.930000    -121.800000
50%         3.534800      29.000000 ... 34.260000    -118.490000
75%         4.743250      37.000000 ... 37.710000    -118.010000
max        15.000100      52.000000 ... 41.950000    -114.310000

```

[8 rows x 8 columns]

### Модель без масштабирования данных

```

[51]: california_X_train_1, california_X_test_1, california_y_train_1,
  ↪ california_y_test_1 = train_test_split(
    df_california[california_x_columns], df_california['target'],
    ↪ test_size=0.2, random_state=1)

```

```
california_X_train_1.shape, california_X_test_1.shape
```

```
[51]: ((16512, 8), (4128, 8))
```

```

[52]: svr_1 = SVR()
svr_1.fit(california_X_train_1, california_y_train_1)

```

```

[52]: SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,
  ↪ gamma='scale',

```



```
kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```

```
[0]: california_y_pred_1 = svr_1.predict(california_X_test_1)
```

```
[54]: mean_absolute_error(california_y_test_1, california_y_pred_1),  
      ↪ mean_squared_error(california_y_test_1, california_y_pred_1)
```

```
[54]: (0.8591093473587432, 1.320732383526986)
```

### Модель с масштабированием данных

```
[0]: # Масштабирование данных в диапазоне от 0 до 1  
sc1 = MinMaxScaler()  
sc1_data = sc1.fit_transform(df_california[california_x_columns])
```

```
[56]: sc1_data[:2]
```

```
[56]: array([[0.53966842, 0.78431373, 0.0435123 , 0.02046866, 0.00894083,  
            0.00149943, 0.5674814 , 0.21115538],  
            [0.53802706, 0.39215686, 0.03822395, 0.01892926, 0.0672104 ,  
            0.00114074, 0.565356 , 0.21215139]])
```

```
[57]: california_X_train_2, california_X_test_2, california_y_train_2,  
      ↪ california_y_test_2 = train_test_split(  
        sc1_data, df_california['target'], test_size=0.2, random_state=1)  
  
california_X_train_2.shape, california_X_test_2.shape
```

```
[57]: ((16512, 8), (4128, 8))
```

```
[58]: svr_2 = SVR()  
svr_2.fit(california_X_train_2, california_y_train_2)
```

```
[58]: SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,  
      ↪ gamma='scale',  
        kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```

```
[0]: california_y_pred_2 = svr_2.predict(california_X_test_2)
```

```
[60]: mean_absolute_error(california_y_test_2, california_y_pred_2),  
      ↪ mean_squared_error(california_y_test_2, california_y_pred_2)
```

```
[60]: (0.4505817514262831, 0.45299763645897284)
```

### Эксперимент на основе кросс-валидации:

```
[0]: scoring = {'MAE': 'neg_mean_absolute_error',  
               'MSE': 'neg_mean_squared_error',  
               'R2': 'r2'}
```

```
[62]: scores_1 = cross_validate(SVR(),  
                               df_california[california_x_columns],  
                               ↪ df_california['target'], scoring=scoring,  
                               cv=3, return_train_score=True)
```

```
scores_1
```

```
[62]: {'fit_time': array([12.11082315, 12.1236763 , 12.2641027 ]),  
      'score_time': array([2.76550102, 2.76244593, 2.79872417]),  
      'test_MAE': array([-0.88547721, -0.81302601, -0.95274753]),  
      'test_MSE': array([-1.33180311, -1.2491849 , -1.59989289]),  
      'test_R2': array([ 0.01210917, -0.04535864, -0.12118326]),  
      'train_MAE': array([-0.86838827, -0.90060707, -0.83393092]),  
      'train_MSE': array([-1.35236242, -1.40856087, -1.31120754]),  
      'train_R2': array([-0.02754543, -0.00651499, -0.03005051])}
```

```
[63]: # Без масштабирования  
      -np.mean(scores_1['train_MAE']), -np.mean(scores_1['test_MAE']), \  
      -np.mean(scores_1['train_MSE']), -np.mean(scores_1['test_MSE'])
```

```
[63]: (0.8676420830790912, 0.8837502471655386, 1.3573769428582043, 1.  
      ↪ 393626970258224)
```

```
[64]: scores_2 = cross_validate(SVR(),  
                               sc1_data, df_california['target'],  
                               ↪ scoring=scoring,  
                               cv=3, return_train_score=True)  
scores_2
```

```
[64]: {'fit_time': array([11.35171986, 11.65829659, 11.28841114]),  
      'score_time': array([2.4718802 , 2.4750576 , 2.41563725]),  
      'test_MAE': array([-0.51033856, -0.45036369, -0.56060072]),  
      'test_MSE': array([-0.54870328, -0.433136 , -0.61684615]),  
      'test_R2': array([0.59298868, 0.63753767, 0.56772258]),  
      'train_MAE': array([-0.44006452, -0.46199119, -0.43109582]),  
      'train_MSE': array([-0.43456878, -0.46134241, -0.42690641]),  
      'train_R2': array([0.66980806, 0.67033867, 0.66463343])}
```

## 9. Данные метрик машины опорных векторов для сравнения с другими моделями

```
[65]: # С масштабированием  
      -np.mean(scores_2['train_MAE']), -np.mean(scores_2['test_MAE']), \  
      -np.mean(scores_2['train_MSE']), -np.mean(scores_2['test_MSE']), \  
      np.mean(scores_2['train_R2']), np.mean(scores_2['test_R2'])
```

```
[65]: (0.4443838432807796,  
      0.507100989083478,  
      0.44093920269594206,  
      0.5328951447989484,  
      0.6682600547769003,  
      0.599416311806581)
```


Аномалии

```
[66]: anom_cl = OneClassSVM()
anom_cl.fit(california_x.reshape(-1, 1))
np.unique(anom_cl.predict(california_x.reshape(-1, 1)))
```

```
[66]: array([-1,  1])
```

```
[67]: plot_regr(OneClassSVM())
```

<bound method BaseEstimator.\_repr\_ of OneClassSVM(cache\_size=200, coef0=0.0, degree=3, gamma='scale', kernel='rbf', max\_iter=-1, nu=0.5, shrinking=True, tol=0.001, verbose=False)>



Дерево решений для регрессии

```
[0]: def random_dataset_for_regression():
    """
    Создание случайного набора данных для регрессии
    """
    rng = np.random.RandomState(1)
    X_train = np.sort(5 * rng.rand(80, 1), axis=0)
    y_train = np.sin(X_train).ravel()
    y_train[::5] += 3 * (0.5 - rng.rand(16))
    X_test = np.arange(0.0, 5.0, 0.01)[:, np.newaxis]
    return X_train, y_train, X_test
```

```
[0]: def plot_tree_regression(X_train, y_train, X_test):
    """
    Построение деревьев и вывод графиков для заданного датасета
    """

    # Обучение регрессионной модели
    regr_1 = DecisionTreeRegressor(max_depth=3)
    regr_2 = DecisionTreeRegressor(max_depth=10)
    regr_1.fit(X_train, y_train)
    regr_2.fit(X_train, y_train)

    # Предсказание
```

```

y_1 = regr_1.predict(X_test)
y_2 = regr_2.predict(X_test)

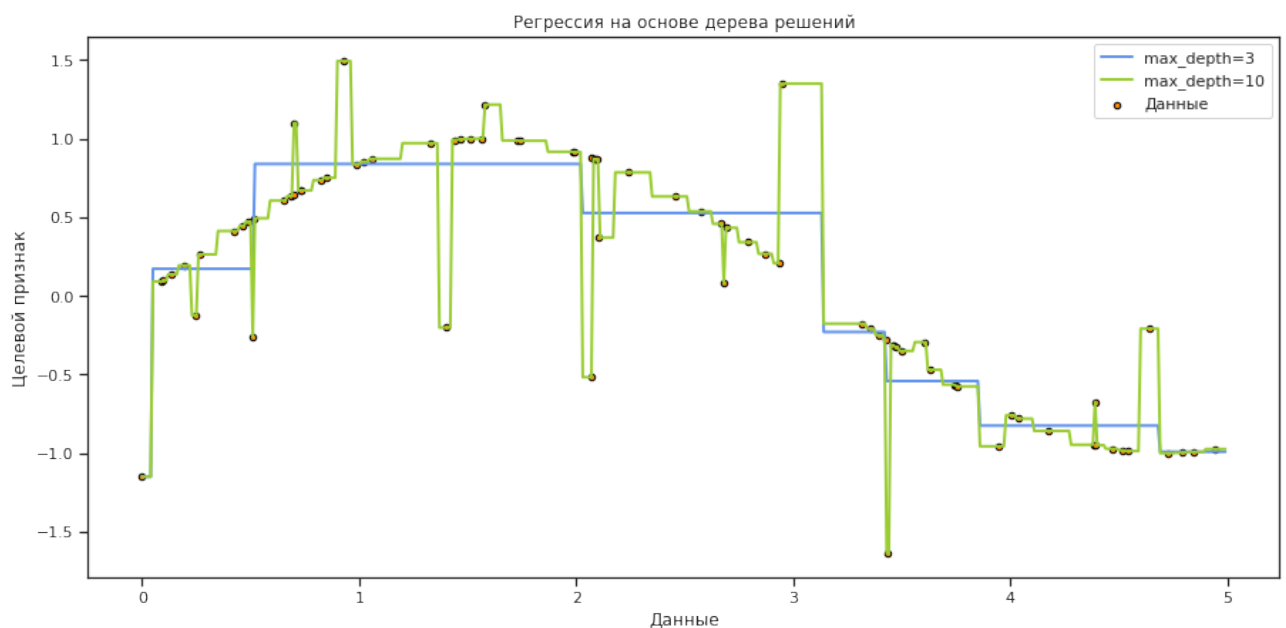
# Вывод графика
fig, ax = plt.subplots(figsize=(15,7))
plt.scatter(X_train, y_train, s=20, edgecolor="black",
↪c="darkorange", label="Данные")
plt.plot(X_test, y_1, color="cornflowerblue", label="max_depth=3",
↪linewidth=2)
plt.plot(X_test, y_2, color="yellowgreen", label="max_depth=10",
↪linewidth=2)
plt.xlabel("Данные")
plt.ylabel("Целевой признак")
plt.title("Регрессия на основе дерева решений")
plt.legend()
plt.show()

```

```

[70]: X_train, y_train, X_test = random_dataset_for_regression()
plot_tree_regression(X_train, y_train, X_test)

```



## 10. Дерево решений

```

[71]: california_tree_regr = DecisionTreeRegressor(random_state=1)
california_tree_regr.fit(df_california, data.target)
california_tree_regr

```

```

[71]: DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
max_features=None, max_leaf_nodes=None,

```

```

        min_impurity_decrease=0.0,
    ↪ min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0,
    ↪ presort='deprecated',
        random_state=1, splitter='best')

```

```

[0]: # Визуализация дерева
def get_png_tree(tree_model_param, feature_names_param):
    dot_data = StringIO()
    export_graphviz(tree_model_param, out_file=dot_data,
    ↪ feature_names=feature_names_param,
        filled=True, rounded=True, special_characters=True)
    graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
    return graph.create_png()

```

```

[73]: Image(get_png_tree(california_tree_regr, df_california.columns),
    ↪ height="500")

```

dot: graph is too large for cairo-renderer bitmaps. Scaling by 0.0643085  
 ↪ to fit

[73]:



```

[74]: california_tree_regr_prun = DecisionTreeRegressor(random_state=1,
    ↪ max_depth=3)
    california_tree_regr_prun.fit(df_california, data.target)
    california_tree_regr_prun

```

```

[74]: DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=3,
        max_features=None, max_leaf_nodes=None,
        min_impurity_decrease=0.0,
    ↪ min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0,
    ↪ presort='deprecated',
        random_state=1, splitter='best')

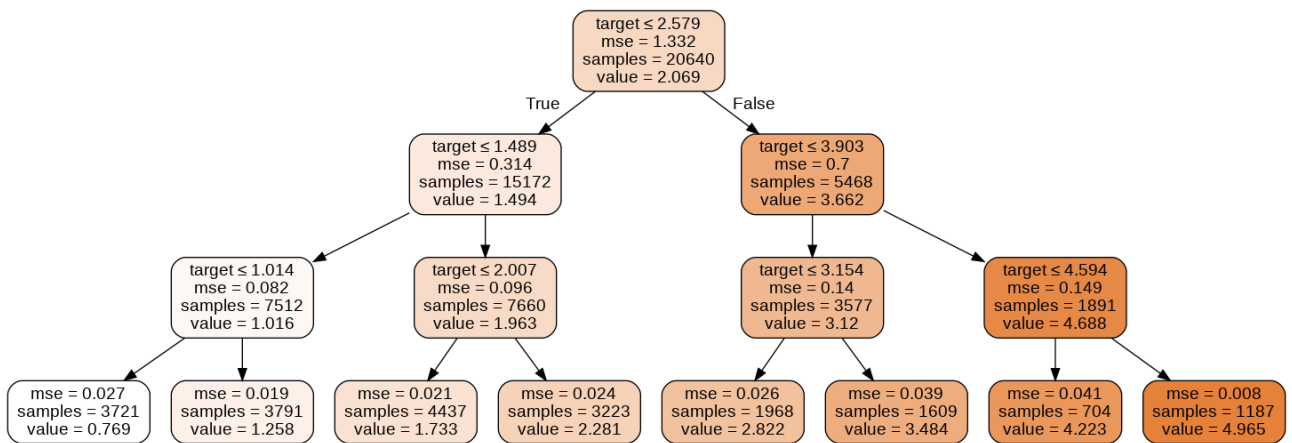
```

```

[75]: Image(get_png_tree(california_tree_regr_prun, df_california.columns),
    ↪ height="500")

```

[75]:



```
[76]: scores_4 = cross_validate(DecisionTreeRegressor(random_state=1,
    ↪max_depth=3),
                                sc1_data, df_california['target'],
    ↪scoring=scoring,
                                cv=3, return_train_score=True)
scores_4
```

```
[76]: {'fit_time': array([0.05726457, 0.04353571, 0.03309608]),
'score_time': array([0.00187492, 0.00159335, 0.00155783]),
'test_MAE': array([-0.60197332, -0.5610655 , -0.65599163]),
'test_MSE': array([-0.63497084, -0.5636626 , -0.78360926]),
'test_R2': array([0.52899805, 0.52830877, 0.45085724]),
'train_MAE': array([-0.59865069, -0.60456721, -0.55194598]),
'train_MSE': array([-0.61718838, -0.64270341, -0.56133664]),
'train_R2': array([0.53105093, 0.54074358, 0.55902855])}
```

## 11. Данные метрик дерева решений для сравнения с другими моделями

```
[77]: -np.mean(scores_4['train_MAE']), -np.mean(scores_4['test_MAE']), \
      -np.mean(scores_4['train_MSE']), -np.mean(scores_4['test_MSE']), \
      np.mean(scores_4['train_R2']), np.mean(scores_4['test_R2'])
```

```
[77]: (0.5850546241714393,
0.6063434807383564,
0.6070761448129344,
0.6607475666049402,
0.5436076857580303,
0.5027213535306057)
```

## 12. Все методы показали примерно одинаковые результаты, однако в целом наилучший результат дала машина опорных векторов

```
[99]: -np.mean(scores_3['train_MAE']), -np.mean(scores_3['test_MAE']), \
      -np.mean(scores_3['train_MSE']), -np.mean(scores_3['test_MSE']), \
      np.mean(scores_3['train_R2']), np.mean(scores_3['test_R2'])
```

```
[99]: (0.5276885217738251,
      0.5502350422781422,
      0.5178224843539592,
      0.5611078368165351,
      0.6103948517000314,
      0.5762820158960761)
```

```
[79]: -np.mean(scores_2['train_MAE']), -np.mean(scores_2['test_MAE']), \
      -np.mean(scores_2['train_MSE']), -np.mean(scores_2['test_MSE']), \
      np.mean(scores_2['train_R2']), np.mean(scores_2['test_R2'])
```

```
[79]: (0.4443838432807796,
      0.507100989083478,
      0.44093920269594206,
      0.5328951447989484,
      0.6682600547769003,
      0.599416311806581)
```

```
[80]: -np.mean(scores_4['train_MAE']), -np.mean(scores_4['test_MAE']), \
      -np.mean(scores_4['train_MSE']), -np.mean(scores_4['test_MSE']), \
      np.mean(scores_4['train_R2']), np.mean(scores_4['test_R2'])
```

```
[80]: (0.5850546241714393,
      0.6063434807383564,
      0.6070761448129344,
      0.6607475666049402,
      0.5436076857580303,
      0.5027213535306057)
```

## 13. Подбор одного гиперпараметра

### Линейная регрессия

```
[121]: n_range = np.array(range(0,100,5))
      tuned_parameters = [{'n_jobs': n_range}]
      tuned_parameters
```

```
[121]: [{'n_jobs': array([ 0,  5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60,
      ↪ 65, 70,
      75, 80,
      85, 90, 95])}]
```

```
[0]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.model_selection import learning_curve, validation_curve
```

```
[122]: from sklearn.model_selection import KFold, RepeatedKFold, LeaveOneOut,
↳ LeavePOut, ShuffleSplit, StratifiedKFold
reg_gs = GridSearchCV(LinearRegression(), tuned_parameters,
↳ cv=ShuffleSplit(), scoring='neg_mean_absolute_error')
reg_gs.fit(data.data, data.target)
```

```
[122]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=None,
↳ test_size=None,
train_size=None),
error_score=nan,
estimator=LinearRegression(copy_X=True, fit_intercept=True,
n_jobs=None, normalize=False),
iid='deprecated', n_jobs=None,
param_grid=[{'n_jobs': array([ 0, 5, 10, 15, 20, 25, 30,
↳ 35, 40,
45, 50, 55, 60, 65, 70, 75, 80,
85, 90, 95])}],
pre_dispatch='2*n_jobs', refit=True,
↳ return_train_score=False,
scoring='neg_mean_absolute_error', verbose=0)
```

```
[96]: reg_gs.best_estimator_
```

```
[96]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=0,
↳ normalize=False)
```

```
[89]: LinearRegression().get_params().keys()
```

```
[89]: dict_keys(['copy_X', 'fit_intercept', 'n_jobs', 'normalize'])
```

```
[98]: scores_lin_improved = cross_validate(reg_gs.best_estimator_,
sc1_data, df_california['target'],
↳ scoring=scoring,
cv=3, return_train_score=True)
scores_lin_improved
```

```
[98]: {'fit_time': array([0.01382542, 0.00902247, 0.00635958]),
'score_time': array([0.00205398, 0.00235367, 0.0018456 ]),
'test_MAE': array([-0.57887213, -0.50877351, -0.56305949]),
'test_MSE': array([-0.59988822, -0.49188049, -0.59155481]),
'test_R2': array([0.55502126, 0.58837838, 0.58544641]),
'train_MAE': array([-0.51394065, -0.54832542, -0.52079949]),
'train_MSE': array([-0.49933994, -0.54483711, -0.50929041]),
'train_R2': array([0.62059396, 0.61067588, 0.59991471])}
```

```
[101]:
```



```
-np.mean(scores_lin_improved['train_MAE']), -np.
    ↳mean(scores_lin_improved['test_MAE']), \
-np.mean(scores_lin_improved['train_MSE']), -np.
    ↳mean(scores_lin_improved['test_MSE']), \
np.mean(scores_lin_improved['train_R2']), np.
    ↳mean(scores_lin_improved['test_R2']))
```

```
[101]: (0.5276885217738251,
        0.5502350422781422,
        0.5178224843539592,
        0.5611078368165351,
        0.6103948517000314,
        0.5762820158960761)
```

Вспоминаем метрики при случайно подобранном гиперпараметре:

```
[100]: -np.mean(scores_3['train_MAE']), -np.mean(scores_3['test_MAE']), \
        -np.mean(scores_3['train_MSE']), -np.mean(scores_3['test_MSE']), \
        np.mean(scores_3['train_R2']), np.mean(scores_3['test_R2']))
```

```
[100]: (0.5276885217738251,
        0.5502350422781422,
        0.5178224843539592,
        0.5611078368165351,
        0.6103948517000314,
        0.5762820158960761)
```

**Для линейной регрессии отличий нет  
Машина опорных векторов**

```
[102]: SVR().get_params().keys()
```

```
[102]: dict_keys(['C', 'cache_size', 'coef0', 'degree', 'epsilon', 'gamma',
    ↳'kernel',
    'max_iter', 'shrinking', 'tol', 'verbose'])
```

```
[103]: n_range = np.array(range(1,20,1))
        tuned_parameters = [{ 'C': n_range}]
        tuned_parameters
```

```
[103]: [{ 'C': array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14,
    ↳15, 16,
        17,
        18, 19])}]
```

```
[105]: reg_gs_svr = GridSearchCV(SVR(), tuned_parameters, cv=3,
    ↳scoring='neg_mean_absolute_error')
        reg_gs_svr.fit(data.data, data.target)
```

```
[105]: GridSearchCV(cv=3, error_score=nan,
    ↳estimator=SVR(C=1.0, cache_size=200, coef0=0.0, degree=3,
    ↳epsilon=0.1, gamma='scale', kernel='rbf',
    ↳max_iter=-1, shrinking=True, tol=0.001,
```

```

        verbose=False),
        iid='deprecated', n_jobs=None,
        param_grid=[{'C': array([ 1,  2,  3,  4,  5,  6,  7,  8, 9,
→ 9, 10,
11, 12, 13, 14, 15, 16, 17,
        18, 19])}]},
        pre_dispatch='2*n_jobs', refit=True,
→ return_train_score=False,
        scoring='neg_mean_absolute_error', verbose=0)

```

```
[106]: reg_gs_svr.best_estimator_
```

```
[106]: SVR(C=19, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,
→ gamma='scale',
        kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)

```

```
[107]: scores_svr_improved = cross_validate(reg_gs_svr.best_estimator_,
        sc1_data, df_california['target'],
→ scoring=scoring,
        cv=3, return_train_score=True)
scores_svr_improved

```

```
[107]: {'fit_time': array([15.53939581, 16.61715221, 15.49336576]),
'score_time': array([2.32830477, 2.38654304, 2.32462382]),
'test_MAE': array([-0.49278758, -0.4468058 , -0.541635  ]),
'test_MSE': array([-0.52777176, -0.40829368, -0.63635923]),
'test_R2': array([0.60851505, 0.65832654, 0.55404807]),
'train_MAE': array([-0.41046405, -0.42600974, -0.40093578]),
'train_MSE': array([-0.39641595, -0.41509394, -0.38038344]),
'train_R2': array([0.69879716, 0.70338643, 0.70118067])}

```

```
[108]: -np.mean(scores_svr_improved['train_MAE']), -np.
→ mean(scores_svr_improved['test_MAE']), \
-np.mean(scores_svr_improved['train_MSE']), -np.
→ mean(scores_svr_improved['test_MSE']), \
np.mean(scores_svr_improved['train_R2']), np.
→ mean(scores_svr_improved['test_R2'])

```

```
[108]: (0.4124698590661901,
0.49374279450566805,
0.397297772757442,
0.5241415550388026,
0.7011214195108509,
0.6069632218804634)

```

Данные с произвольным параметром

```
[109]: -np.mean(scores_2['train_MAE']), -np.mean(scores_2['test_MAE']), \
-np.mean(scores_2['train_MSE']), -np.mean(scores_2['test_MSE']), \
np.mean(scores_2['train_R2']), np.mean(scores_2['test_R2'])

```

```
[109]: (0.4443838432807796,  
        0.507100989083478,  
        0.44093920269594206,  
        0.5328951447989484,  
        0.6682600547769003,  
        0.599416311806581)
```

**В целом результат стал лучше для машины опорных векторов**  
**Решающее дерево**

```
[110]: DecisionTreeRegressor().get_params().keys()
```

```
[110]: dict_keys(['ccp_alpha', 'criterion', 'max_depth', 'max_features',  
                'max_leaf_nodes', 'min_impurity_decrease', 'min_impurity_split',  
                'min_samples_leaf', 'min_samples_split', 'min_weight_fraction_leaf',  
                ↪ 'presort',  
                'random_state', 'splitter'])
```

```
[111]: n_range = np.array(range(1,7,1))  
tuned_parameters = [{'max_depth': n_range}]  
tuned_parameters
```

```
[111]: [{'max_depth': array([1, 2, 3, 4, 5, 6])}]
```

```
[118]: reg_gs_dt = GridSearchCV(DecisionTreeRegressor(random_state=1),  
    ↪ tuned_parameters,  
                                     cv=3, scoring='neg_mean_absolute_error')  
reg_gs_dt.fit(data.data, data.target)
```

```
[118]: GridSearchCV(cv=3, error_score=nan,  
                  estimator=DecisionTreeRegressor(ccp_alpha=0.0,  
    ↪ criterion='mse',  
                                                    max_depth=None,  
    ↪ max_features=None,  
                                                    max_leaf_nodes=None,  
                                                    min_impurity_decrease=0.0,  
                                                    min_impurity_split=None,  
                                                    min_samples_leaf=1,  
                                                    min_samples_split=2,  
                                                    min_weight_fraction_leaf=0.  
    ↪ 0,  
                                                    presort='deprecated',  
                                                    random_state=1,  
    ↪ splitter='best'),  
                  iid='deprecated', n_jobs=None,  
                  param_grid=[{'max_depth': array([1, 2, 3, 4, 5, 6])}],  
                  pre_dispatch='2*n_jobs', refit=True,  
    ↪ return_train_score=False,  
                  scoring='neg_mean_absolute_error', verbose=0)
```

```
[120]: reg_gs_dt.best_estimator_
```

```
[120]: DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=6,
                             max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0,
                             ↪min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0,
                             ↪presort='deprecated',
                             random_state=1, splitter='best')
```

```
[123]: scores_dt_improved = cross_validate(reg_gs_dt.best_estimator_,
                                             sc1_data, df_california['target'],
                                             ↪scoring=scoring,
                                             cv=3, return_train_score=True)
scores_dt_improved
```

```
[123]: {'fit_time': array([0.07014084, 0.06682444, 0.06235528]),
        'score_time': array([0.00182557, 0.00171566, 0.00167012]),
        'test_MAE': array([-0.61119449, -0.49988436, -0.59795421]),
        'test_MSE': array([-0.69455089, -0.47155863, -0.70879242]),
        'test_R2': array([0.48480339, 0.60538437, 0.50328787]),
        'train_MAE': array([-0.45610619, -0.48600847, -0.40178872]),
        'train_MSE': array([-0.40015903, -0.44928409, -0.32800898]),
        'train_R2': array([0.69595312, 0.67895518, 0.74232468])}
```

```
[124]: -np.mean(scores_dt_improved['train_MAE']), -np.
        ↪mean(scores_dt_improved['test_MAE']), \
-np.mean(scores_dt_improved['train_MSE']), -np.
        ↪mean(scores_dt_improved['test_MSE']), \
np.mean(scores_dt_improved['train_R2']), np.
        ↪mean(scores_dt_improved['test_R2'])
```

```
[124]: (0.4479677942838462,
        0.5696776857508322,
        0.39248403231981016,
        0.624967314819766,
        0.7057443261913532,
        0.5311585429990194)
```

Метрики для произвольного параметра:

```
[125]: -np.mean(scores_4['train_MAE']), -np.mean(scores_4['test_MAE']), \
        -np.mean(scores_4['train_MSE']), -np.mean(scores_4['test_MSE']), \
        np.mean(scores_4['train_R2']), np.mean(scores_4['test_R2'])
```

```
[125]: (0.5850546241714393,
        0.6063434807383564,
        0.6070761448129344,
        0.6607475666049402,
        0.5436076857580303,
        0.5027213535306057)
```

В целом результаты для дерева решений стали лучше

## 14. Итого:

```
[126]: -np.mean(scores_lin_improved['train_MAE']), -np.  
        ↪mean(scores_lin_improved['test_MAE']), \  
-np.mean(scores_lin_improved['train_MSE']), -np.  
        ↪mean(scores_lin_improved['test_MSE']), \  
np.mean(scores_lin_improved['train_R2']), np.  
        ↪mean(scores_lin_improved['test_R2'])
```

```
[126]: (0.5276885217738251,  
        0.5502350422781422,  
        0.5178224843539592,  
        0.5611078368165351,  
        0.6103948517000314,  
        0.5762820158960761)
```

```
[127]: -np.mean(scores_svr_improved['train_MAE']), -np.  
        ↪mean(scores_svr_improved['test_MAE']), \  
-np.mean(scores_svr_improved['train_MSE']), -np.  
        ↪mean(scores_svr_improved['test_MSE']), \  
np.mean(scores_svr_improved['train_R2']), np.  
        ↪mean(scores_svr_improved['test_R2'])
```

```
[127]: (0.4124698590661901,  
        0.49374279450566805,  
        0.3972977772757442,  
        0.5241415550388026,  
        0.7011214195108509,  
        0.6069632218804634)
```

```
[128]: -np.mean(scores_dt_improved['train_MAE']), -np.  
        ↪mean(scores_dt_improved['test_MAE']), \  
-np.mean(scores_dt_improved['train_MSE']), -np.  
        ↪mean(scores_dt_improved['test_MSE']), \  
np.mean(scores_dt_improved['train_R2']), np.  
        ↪mean(scores_dt_improved['test_R2'])
```

```
[128]: (0.4479677942838462,  
        0.5696776857508322,  
        0.39248403231981016,  
        0.624967314819766,  
        0.7057443261913532,  
        0.5311585429990194)
```

**#Вывод: #Машина опорных векторов показывает наилучшие результаты, на втором месте решающее дерево, на третьем - линейная регрессия.**